

CS3500 : Operating Systems

$O(1)$  Scheduler for xv6

# Team Sooners

Viswanath Tadi  
CS18B047  
G V S Praveen  
CS18B017

uint64 priority is added  
to store the static priority  
of the process.

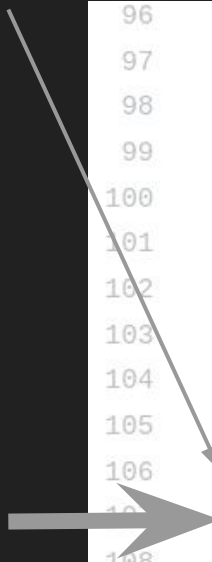
Static Priority initialized  
to 20

```
96
97 // these are private to the process, so p->lock need not be
98 uint64 kstack;           // Bottom of kernel stack for t
99 uint64 sz;               // Size of process memory (byte
100 pagetable_t pagetable;   // Page table
101 struct trapframe *tf;    // data page for trampoline.S
102 struct context context;   // swtch() here to run process
103 struct file *ofile[NOFILE]; // Open files
104 struct inode *cwd;        // Current directory
105 char name[16];           // Process name (debugging)
106 uint64 priority;          // static priority
107 uint64 dynamic_priority;  // dynamic priority of the
108 struct proc* next;        // next process pointer in
109 int last_queue;           // queue from which
110 uint64 running_since;
111 uint64 sleeping_since;
112 uint64 total_runtime;
113 uint64 total_sleeptime;
```

uint64 priority is added to store the static priority of the process.

Static Priority initialized to 20

uint64 dynamic\_priority is added to store the dynamic priority of the process




```
96
97 // these are private to the process, so p->lock need not be
98 uint64 kstack; // Bottom of kernel stack for t
99 uint64 sz; // Size of process memory (byte
100 pagetable_t pagetable; // Page table
101 struct trapframe *tf; // data page for trampoline.S
102 struct context context; // switch() here to run process
103 struct file *ofile[NOFILE]; // Open files
104 struct inode *cwd; // Current directory
105 char name[16]; // Process name (debugging)
106 uint64 priority; // static priority
107 uint64 dynamic_priority; // dynamic priority of the
108 struct proc* next; // next process pointer in
109 int last_queue; // queue from which
110 uint64 running_since;
111 uint64 sleeping_since;
112 uint64 total_runtime;
113 uint64 total_sleeptime;
```

System call `sys_setpriority()` allows a process to change its own static priority.

Meant for use in testing programs

System call `sys_getpriority()` allows a process to read its own dynamic priority.

Meant for use in testing programs



```
98
99  uint64
100  sys_setpriority(void)
101  {
102      int n;
103      if(argint(0, &n) < 0)
104          return -1;
105      struct proc* p = myproc();
106      if(n>39 || n<0) return -1;
107      p->priority=n;
108      return 0;
109  }
110
111  uint64
112  sys_getpriority(void)
113  {
114      struct proc* p = myproc();
115      int n;
116      n = p->dynamic_priority;
117      return n;
118  }
119
```

Added support for **nice** command from shell.

\$ my\_prog -nice 10

- Check if second last argument is **“-nice”**
- Convert last argument to integer
- Update priority of process
- Remove the last two arguments from argv before passing on to the process.

```

80     if(argcval > 2 && strlen(argv[argcval-2])==5 && argv[argcval-2][0]=='-' && argv[argcval-2][1]=='n'
81         && argv[argcval-2][2]=='i' && argv[argcval-2][3]=='c' && argv[argcval-2][4]=='e')
82     {
83         if(strlen(argv[argcval-1])<=2)
84         {
85             if(strlen(argv[argcval-1])==2 && argv[argcval-1][0]>='0' && argv[argcval-1][0]<='9' &&
86                 argv[argcval-1][1]>='0' && argv[argcval-1][1]<='9')
87             {
88                 int val=0;
89                 val+=(argv[argcval-1][0]-'0') * 10;
90                 val+=(argv[argcval-1][1]-'0');
91                 if(val<=39 && val>=0)
92                     p->priority=val;
93             }
94             else if(strlen(argv[argcval-1])==1 && argv[argcval-1][0]>='0' && argv[argcval-1][0]<='9')
95             {
96                 int val=0;
97                 val+=(argv[argcval-1][0]-'0');
98                 if(val<=39 && val>=0)
99                     p->priority=val;
100             }
101         }
102         argcval-=2;
103         argv[argcval] = 0;
104     }

```

struct proc\* q contains two arrays of 40 pointers each. Two arrays act as active and passive queues interchangeably. 40 entries denote the heads of 40 priority queues.

struct proc\* qlast contains two arrays of 40 pointers each. 40 entries denote the tails of 40 priority queues.

Denotes which of q[0] and q[1] is active


Necessary spinlocks for the above structures. Help synchronize access across cores.

```
9 //Linked list implementation
10
11 struct sched_queue
12 {
13     struct proc* q[2][40];
14     struct proc* qlast[2][40];
15     int sched_active;
16     struct spinlock lock[2][40];
17     struct spinlock lock_active;
18 };
19
```

Initialize active and passive queues to empty.

Initialize q[0] as the active queue by setting sched\_active to 0


```
22 void
23 sched_init()
24 {
25     int i;
26     for(i=0;i<40;i++)
27     {
28         queue.q[0][i]=queue.q[1][i]=queue.qlast[0][i]=queue.qlast[1][i]=0;
29     }
30     queue.sched_active = 0;
31 }
32
```





`sched_init()` called from `main.c` before calling scheduler.

`sched_init()` initializes the active and passive queues to all empty queues.



```
27     binit();           // buffer cache
28     iinit();           // inode cache
29     fileinit();        // file table
30     virtio_disk_init(); // emulated hard disk
31     sched_init();       // initializing scheduler structures
32     userinit();         // first user process
33     __sync_synchronize();
34     started = 1;
35 } else {
36     while(started == 0)
37         ;
38     __sync_synchronize();
39     printf("hart %d starting\n", cpuid());
40     hminithart();       // turn on timer
```

Dynamic priority of a process is calculated just before it is inserted into a queue.

Bonus = (sleepime\*11)/  
(sleepime + runtime)

Dynamic priority =  
Static\_priority - bonus + 5

Dynamic priority in the  
struct proc is updated just  
before insertion.

```

33 void
34 sched_insert(struct proc* curp,int active)
35 {
36     if(active!=0 && active!=1)panic("active value out of bounds");
37     curp->next = 0;
38     int bonus = 5;
39     if(!(curp->total_sleepime==0 && curp->total_runtime==0)){
40         bonus = ((curp->total_sleepime*11)/(curp->total_sleepime+curp->total_runtime));
41     }
42     int dprio = curp->priority - bonus + 5;
43     if(dprio <= 0) dprio = 0;
44     if(dprio >= 39) dprio = 39;
45     curp->dynamic_priority = dprio;
46     int p = curp->dynamic_priority;

```

⋮

If active is 1, process is inserted into currently active queue. If active is 0, process is inserted into the queue which was passive when it was allocated to the CPU. Index denotes the queue into which insertion takes place.

Here, p is the dynamic priority.  
Process curp is inserted into queue denoted by index and priority queue denoted by p.  
Insertion is done at the end of the queue in O(1) time bounds.

Necessary locking is present in place to ensure synchronization.

kernel / sched.c / sched\_insert()

```
47     int index;
48
49     acquire(&queue.lock_active);
50     if(active == 1)
51     {
52         index = queue.sched_active;
53     }
54     else if (active == 0)
55     {
56         index = 1 - curp->last_queue;
57     }
58     acquire(&queue.lock[index][p]);
59     if(queue.q[index][p]==0)
60     {
61         queue.q[index][p]=curp;
62         queue.qlast[index][p]=curp;
63     }
64     else
65     {
66         queue.qlast[index][p]->next=curp;
67         queue.qlast[index][p]=curp;
68     }
69     release(&queue.lock_active);
70     release(&queue.lock[index][p]);
71 }
72
```

Iterate through all the active queues in decreasing order of priority to find the first non-empty queue.

If found, remove the first process from queue and return it.

If not found, make the other queue as active and vice versa by changing sched\_active.

```

73  struct proc*
74  sched_get()
75  {
76      int i;
77      struct proc* p;
78      sched_L:
79      acquire(&queue.lock_active);
80      for(i=0;i<40;i++)
81      {
82          acquire(&queue.lock[queue.sched_active][i]);
83          if(queue.q[queue.sched_active][i]!=0)
84          {
85              p = queue.q[queue.sched_active][i];
86              p->last_queue = queue.sched_active;
87              queue.q[queue.sched_active][i] = queue.q[queue.sched_active][i]->next;
88              if(p->next==0)
89              {
90                  queue.qlast[queue.sched_active][i]=0;
91              }
92              release(&queue.lock[queue.sched_active][i]);
93              release(&queue.lock_active);
94              return p;
95          }
96          release(&queue.lock[queue.sched_active][i]);
97      }
98      queue.sched_active = 1 - queue.sched_active;

```

Iterate through all the new active queues in decreasing order of priority to find the first non-empty queue.

If found, remove the first process from queue and return it.

If not found, run the whole function again with the other queue as active.

```


100     for(i=0;i<40;i++)
101     {
102         acquire(&queue.lock[queue.sched_active][i]);
103         if(queue.q[queue.sched_active][i]!=0)
104         {
105             p = queue.q[queue.sched_active][i];
106             p->last_queue = queue.sched_active;
107             queue.q[queue.sched_active][i] = queue.q[queue.sched_active][i]->next;
108             if(p->next==0)
109             {
110                 queue.qlast[queue.sched_active][i]=0;
111             }
112             release(&queue.lock[queue.sched_active][i]);
113             release(&queue.lock_active);
114             return p;
115         }
116         release(&queue.lock[queue.sched_active][i]);
117     }
118
119     queue.sched_active = 1 - queue.sched_active;
120     release(&queue.lock_active);
121     goto sched_L;
122 }

```

When the process is created, following values are initialized

- static priority to 20.
- total\_runtime to 0.
- total\_sleeptime to 0.


```
116     p->pagetable = proc_pagetable(p);
117
118     // Set up new context to start executing at forkret,
119     // which returns to user space.
120     memset(&p->context, 0, sizeof p->context);
121     p->context.ra = (uint64)forkret;
122     p->context.sp = p->kstack + PGSIZE;
123
124     // Make static priority = 5
125     p->priority = 20;
126     p->total_runtime = 0;
127     p->total_sleeptime = 0;
128
129     return p;
130 }
131
```



When the first user process becomes RUNNABLE, insert it into the active queue.

```
217     p->tf->sp = PGSIZE; // user stack pointer
218
219     safestrcpy(p->name, "initcode", sizeof(p->name));
220     p->cwd = namei("/");
221
222     p->state = RUNNABLE;
223
224     sched_insert(p,1);
225
226     release(&p->lock);
227 }
228
229 // Grow or shrink user memory by n bytes.
```

When the child process becomes RUNNABLE, insert it into the active queue.



```
285     safestrcpy(np->name, p->name, sizeof(p->name));
286
287     pid = np->pid;
288
289     np->state = RUNNABLE;
290
291     sched_insert(np,1);
292
293     release(&np->lock);
294
295     return pid;
296 }
297
298 // Pass p's abandoned children to init.
```



Find the next process to schedule using sched\_get().

Calculate time-slice using the dynamic priority of the process.


Configure timer interrupt by updating CLINT\_MTIMECMP to CLINT\_MTIME + timeslice

Update p->running\_since to CLINT\_MTIME since process is set to RUNNING state.


Switch to process's context.

```
459     for(;;){
460         // Avoid deadlock by ensuring that devices can interrupt.
461         intr_on();
462
463         // get next process
464         p = sched_get();
465
466         // timeslice calculation
467         int timeslice;
468         if(p->dynamic_priority < 20)
469             timeslice = (40-p->dynamic_priority)*500000;
470         else
471             timeslice = (40-p->dynamic_priority)*250000;
472
473         *(uint64*)CLINT_MTIMECMP(cpuid()) = *(uint64*)CLINT_MTIME + timeslice;
474
475         acquire(&p->lock);
476
477         // Switch to chosen process. It is the process's job
478         // to release its lock and then reacquire it
479         // before jumping back to us.
480         // printf("%s,%d\n",p->name,p->dynamic_priority);
481         p->running_since = *(uint64*)CLINT_MTIME;
482         p->state = RUNNING;
483         c->proc = p;
484         switch(&c->scheduler, &p->context);
```

Calculate run time using current MTIME and p->running\_since, and add it to p->total\_runtime.




Insert it into the passive queue since it's state is now RUNNABLE.



```
520 // Give up the CPU for one scheduling round.
521 void
522 yield(void)
523 {
524     struct proc *p = myproc();
525     acquire(&p->lock);
526     p->total_runtime += *(uint64*)CLINT_MTIME - p->running_since;
527     //printf("%s,%d,%d\n",p->name,*(uint64*)CLINT_MTIME,p->running_since);
528     p->state = RUNNABLE;
529     sched_insert(p,0);
530     sched();
531     release(&p->lock);
532 }
533
```

Calculate run time  
using current MTIME  
and p->running\_since,  
and add it to  
p->total\_runtime.




```
377 // Parent might be sleeping in wait().
378 wakeup1(original_parent);
379
380 p->total_runtime += *(uint64*)CLINT_MTIME - p->running_since;
381 p->xstate = status;
382 p->state = ZOMBIE;
383
384 release(&original_parent->lock);
385
386 // Jump into the scheduler, never to return.
387 sched();
388 panic("zombie exit");
389 }
390
```

Set p->sleeping\_since to current MTIME as state is changed to SLEEPING.


Calculate run time using current MTIME and p->running\_since, and add it to p->total\_runtime.

```
568     if(lk != &p->lock){ //DOC: sleeplock0
569         acquire(&p->lock); //DOC: sleeplock1
570         release(lk);
571     }
572
573     // Go to sleep.
574     p->total_runtime += *(uint64*)CLINT_MTIME - p->running_since;
575     p->sleeping_since = *(uint64*)CLINT_MTIME;
576     p->chan = chan;
577     p->state = SLEEPING;
578
579     sched();
580
581     // Tidy up
```

Calculate sleep time using current MTIME and p->sleeping\_since, and add it to p->total\_sleeptime.




Since the process becomes RUNNABLE, insert it into the active queue.




```
593 void
594 wakeup(void *chan)
595 {
596     struct proc *p;
597
598     for(p = proc; p < &proc[NPROC]; p++) {
599         acquire(&p->lock);
600         if(p->state == SLEEPING && p->chan == chan) {
601             p->total_sleeptime += *(uint64*)CLINT_MTIME - p->sleeping_since;
602             p->state = RUNNABLE;
603             sched_insert(p,1);
604         }
605         release(&p->lock);
606     }
607 }
608
```

Similar changes made in wakeup1.

Calculate sleep time using current MTIME and p->sleeping\_since, and add it to p->total\_sleeptime.



Since the process becomes RUNNABLE, insert it into the active queue.



```
632     acquire(&p->lock);
633     if(p->pid == pid){
634         p->killed = 1;
635         if(p->state == SLEEPING){
636             // Wake process from sleep().
637             p->total_sleeptime += *(uint64*)CLINT_MTIME - p->sleeping_since;
638             p->state = RUNNABLE;
639             sched_insert(p,1);
640         }
641         release(&p->lock);
642         return 0;
643     }
644     release(&p->lock);
645 }
646 return -1;
```